

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
```

```
In [2]: df= pd.read_csv('Salary_dataset.csv')
```


```
In [3]: df.head()
```

```
Out[3]:
```

	Unnamed: 0	YearsExperience	Salary
0	0	1.2	39344.0
1	1	1.4	46206.0
2	2	1.6	37732.0
3	3	2.1	43526.0
4	4	2.3	39892.0

```
In [4]: # Scatter Plot
plt.scatter(df['YearsExperience'],df['Salary'])
plt.xlabel("YearsExperience")
plt.ylabel("Salary")
```

```
Out[4]: Text(0, 0.5, 'Salary')
```



```
In [5]: # Correlation
df.corr()
```

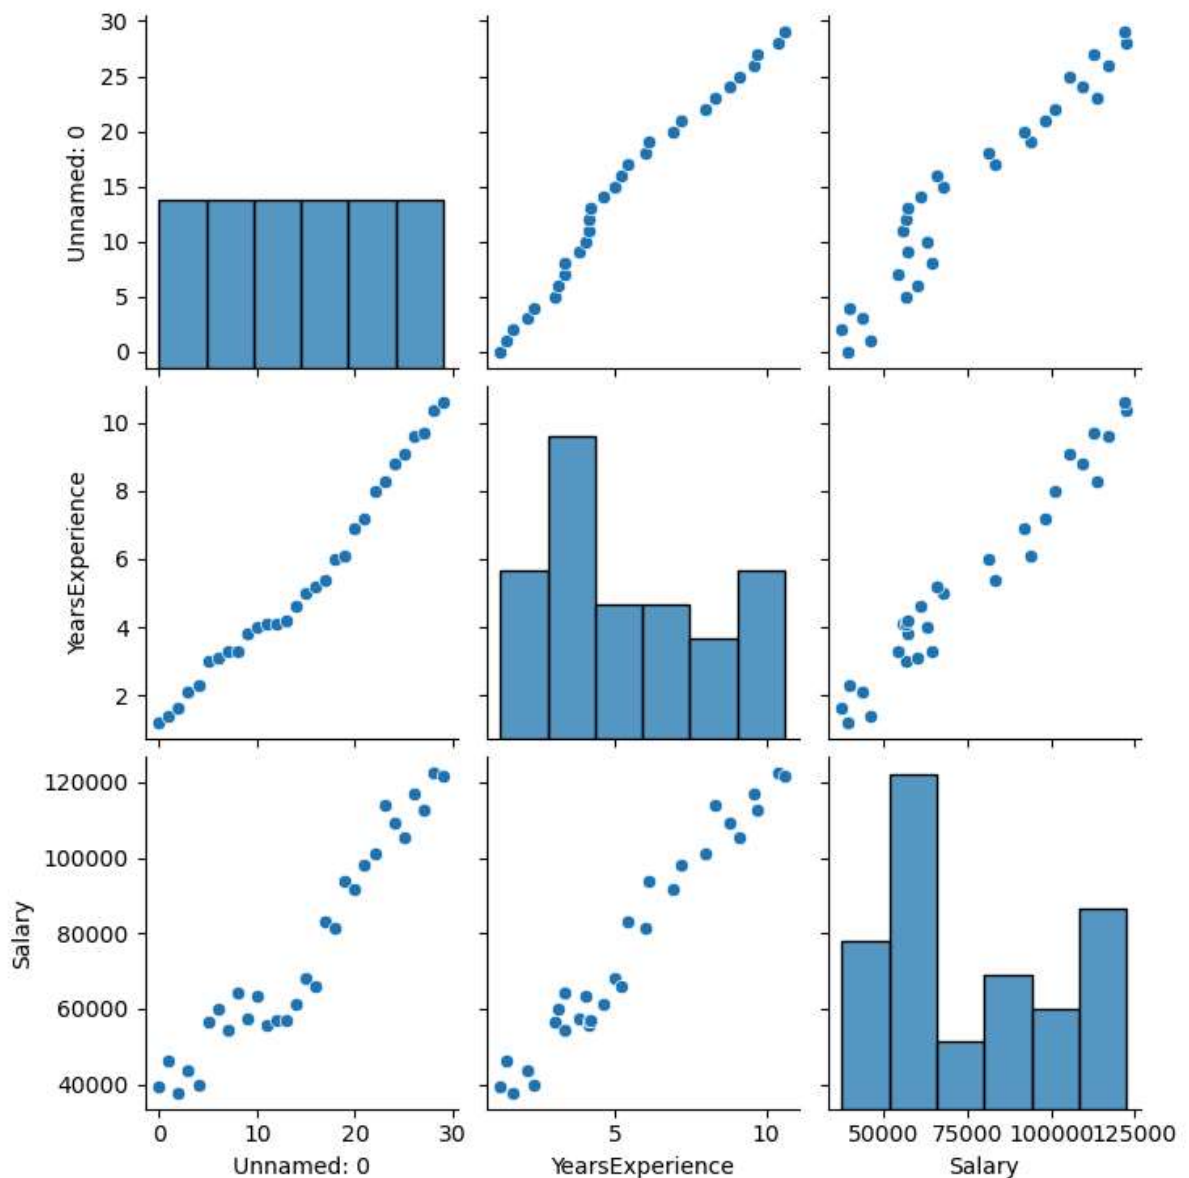
```
Out[5]:
```

	Unnamed: 0	YearsExperience	Salary
Unnamed: 0	1.000000	0.986460	0.960826
YearsExperience	0.986460	1.000000	0.978242
Salary	0.960826	0.978242	1.000000

```
In [7]: # Seaborn for Visualization Plot
import seaborn as sns
sns.pairplot(df)
```

C:\Users\win 10\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
 self._figure.tight_layout(*args, **kwargs)

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x1ef59dc3210>
```



```
In [8]: # Independent and dependent feature
df.head()
```

```
Out[8]:
```

	Unnamed: 0	YearsExperience	Salary
0	0	1.2	39344.0
1	1	1.4	46206.0
2	2	1.6	37732.0
3	3	2.1	43526.0
4	4	2.3	39892.0

```
In [11]: X=df[["YearsExperience"]] ## independent features should be data frame or 2 dim
np.array(X)
```

```
Out[11]: array([[ 1.2],
 [ 1.4],
 [ 1.6],
 [ 2.1],
 [ 2.3],
 [ 3. ],
 [ 3.1],
 [ 3.3],
 [ 3.3],
 [ 3.8],
 [ 4. ],
 [ 4.1],
 [ 4.1],
 [ 4.2],
 [ 4.6],
 [ 5. ],
 [ 5.2],
 [ 5.4],
 [ 6. ],
 [ 6.1],
 [ 6.9],
 [ 7.2],
 [ 8. ],
 [ 8.3],
 [ 8.8],
 [ 9.1],
 [ 9.6],
 [ 9.7],
 [10.4],
 [10.6]])
```

```
In [12]: X=df[["YearsExperience"]]
np.array(X).shape
```

```
Out[12]: (30, 1)
```

```
In [13]: X_series=df["YearsExperience"]  
np.array(X_series).shape
```

```
Out[13]: (30,)
```

```
In [14]: X=df[["YearsExperience"]]  
y=df["Salary"] # this variable can be in series or 1d array
```

```
In [15]: y
```

```
Out[15]: 0      39344.0  
1      46206.0  
2      37732.0  
3      43526.0  
4      39892.0  
5      56643.0  
6      60151.0  
7      54446.0  
8      64446.0  
9      57190.0  
10     63219.0  
11     55795.0  
12     56958.0  
13     57082.0  
14     61112.0  
15     67939.0  
16     66030.0  
17     83089.0  
18     81364.0  
19     93941.0  
20     91739.0  
21     98274.0  
22    101303.0  
23    113813.0  
24    109432.0  
25    105583.0  
26    116970.0  
27    112636.0  
28    122392.0  
29    121873.0  
Name: Salary, dtype: float64
```

```
In [16]: np.array(y).shape
```

```
Out[16]: (30,)
```

```
In [18]: # Train Test Split  
from sklearn.model_selection import train_test_split
```

```
In [19]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=
```

```
In [20]: X_train.shape
```

```
Out[20]: (22, 1)
```

```
In [21]: # Standardization  
from sklearn.preprocessing import StandardScaler
```

```
In [23]: scaler=StandardScaler()  
scaler.fit_transform(X_train)
```

```
Out[23]: array([[ -0.29741739],  
                [-1.38171623],  
                [-0.97043047],  
                [ 0.11386837],  
                [-0.70870316],  
                [-0.26002778],  
                [-0.29741739],  
                [ 1.1607776 ],  
                [-1.306937  ],  
                [-1.23215777],  
                [ 1.57206337],  
                [-1.0452097 ],  
                [ 0.86166068],  
                [ 1.75901144],  
                [ 0.41298529],  
                [ 2.13290759],  
                [ 0.74949183],  
                [-0.59653431],  
                [-0.33480701],  
                [-0.11046932],  
                [ 0.45037491],  
                [-0.67131355]])
```

```
In [24]: X_test=scaler.transform(X_test)
```

```
In [25]: X_test
```

```
Out[25]: array([[ 1.79640106],  
                [ 0.03908914],  
                [ 1.27294644],  
                [ 0.1886476 ],  
                [-0.59653431],  
                [-0.40958624],  
                [ 2.05812836],  
                [ 1.45989452]])
```

```
In [26]: ## Apply Simple Linear Regression  
from sklearn.linear_model import LinearRegression
```

```
In [27]: regression=LinearRegression(n_jobs=-1)
```

```
In [28]: regression.fit(X_train,y_train)
```

```
Out[28]: LinearRegression(n_jobs=-1)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [30]: regression.coef_
```

```
Out[30]: array([9371.0160797])
```

```
In [31]: regression.intercept_
```

```
Out[31]: 24542.025828030746
```


```
In [32]: print("Coefficient or slope:",regression.coef_)  
print("Intercept:",regression.intercept_)
```

Coefficient or slope: [9371.0160797]

Intercept: 24542.025828030746

```
In [33]: ## plot Training data plot best fit line  
plt.scatter(X_train,y_train)  
plt.plot(X_train,regression.predict(X_train))
```

```
Out[33]: [<matplotlib.lines.Line2D at 0x1ef5dd2ea50>]
```



prediction of test data

**predicted height output= intercept
+coef_(YearsExperience)**

y_pred_test =24542.02 + 9371.01(X_test

```
In [34]: ## Predict for test data
y_pred=regression.predict(X_test)
```

C:\Users\win 10\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning:
X does not have valid feature names, but LinearRegression was fitted with fea
ture names
warnings.warn(

```
In [35]: ## Performance Metrics
from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
In [36]: mse=mean_squared_error(y_test,y_pred)
mae=mean_absolute_error(y_test,y_pred)
rms= np.sqrt(mse)
print(mse)
print(mae)
print(rms)
```

```
3847398784.0427666
60020.585847156995
62027.40349267222
```

R square

Formula

$$R^2 = 1 - SSR/SST$$

**R^2 = coefficient of determination SSR = sum of
squares of residuals SST = total sum of
squares**

```
In [37]: from sklearn.metrics import r2_score
```

```
In [38]: score=r2_score(y_test,y_pred)
print(score)
```

```
-5.472618128428121
```

$$\text{Adjusted } R^2 = 1 - [(1-R^2)*(n-1)/(n-k-1)]$$

where:

R2: The R2 of the model **n: The number of observations** **k: The number of predictor variables**

```
In [39]: #display adjusted R-squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

```
Out[39]: -6.551387816499474
```

```
In [40]: ## OLS Linear Regression
import statsmodels.api as sm
```

```
In [41]: model=sm.OLS(X_train,y_train).fit()
```

```
In [42]: prediction=model.predict(X_test)
print(prediction)
```

```
[ 1.31707258e-04  2.86591005e-06  9.33289839e-05  1.38311311e-05
 -4.37362795e-05 -3.00297532e-05  1.50896395e-04  1.07035510e-04]
```



```
In [43]: print(model.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:      YearsExperience    R-squared (uncentered):
0.970
Model:              OLS              Adj. R-squared (uncentered):
0.969
Method:             Least Squares    F-statistic:
681.9
Date:               Fri, 30 May 2025  Prob (F-statistic):
1.71e-17
Time:               20:21:42          Log-Likelihood:
-30.415
No. Observations:   22              AIC:
62.83
Df Residuals:       21              BIC:
63.92
Df Model:            1
Covariance Type:    nonrobust
=====
=
                                coef      std err          t      P>|t|      [0.025      0.97
5]
-----
-
Salary      7.332e-05   2.81e-06    26.114    0.000    6.75e-05   7.92e-0
5
=====
=
Omnibus:          0.229   Durbin-Watson:          1.89
6
Prob(Omnibus):    0.892   Jarque-Bera (JB):          0.37
8
Skew:             0.193   Prob(JB):          0.82
8
Kurtosis:         2.487   Cond. No.          1.0
0
=====
=
```

Notes:

[1] R² is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [45]: # Prediction for new data
         regression.predict(scaler.transform([[72]]))
```

```
C:\Users\win 10\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning:
X does not have valid feature names, but StandardScaler was fitted with feature names
```

```
warnings.warn(
C:\Users\win 10\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning:
X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(
```

```
Out[45]: array([259662.04973488])
```